

Paddle Duel RL Arena — Project Report

SUTD 51.512 Reinforcement Learning for Embodied AI, Y2025-2026

Group members:

Name	Student ID	Contributions
Duan Xu	1010728	Designed and implemented Level 5 (two-agent self-play): built the fixed-opponent baseline (5A) and self-play training with an opponent pool (5B) for both the Feature-DQN and CNN-DQN architectures, ran the full-scale training and evaluation, and produced the final competition-ready checkpoints.
Manan Bimal Mehta	1010949	Designed and implemented Levels 1–3 (tabular Q-learning, SARSA, and Expected SARSA agents; exploration-schedule comparison; opponent-zoo robustness evaluation) and Level 4 (Feature-DQN and CNN-DQN pixel-based agents vs. scripted opponents), including all training, evaluation, and baseline comparisons for these levels.
Zheng Chengsheng	1010719	Wrote and verified the final project report, consolidating results and figures across all five levels, cross-checking reported metrics against the actual notebook/training outputs, and compiling the reproducibility documentation.

1. Introduction

This report covers our design, training, and evaluation of reinforcement learning agents for **Paddle Duel**, a simplified Pong-style environment provided for this project. We attempted all five project levels, progressing from tabular control (Level 1-3) through pixel-based function approximation (Level 4) to two-agent self-play (Level 5). Section 9 summarises which levels our final submitted agents support and how performance changes across difficulty.

The environment itself (`paddle_duel_env.py`) was provided and not modified.

Two players (Left, Right) each choose one of three actions per frame (UP, DOWN, STAY); an episode is one rally, ending when a side scores (+10) or after a 60-second/1800-step timeout, in which case cumulative reward decides the winner. Reward also includes a small per-hit bonus (+0.06 for the first 10 hits, +0.02 after) and a tiny movement penalty (-0.001 for UP/DOWN), as

specified in the project instructions.

Reproducibility instructions, the full checkpoint index, and known environment gotchas are documented separately in `README.md` at the project root; this report focuses on design rationale, experiments, and results.

2. Methodology overview

Across all levels we followed the same experimental discipline: agents were trained using a fixed training seed (and, in later levels, an explicit warm-up period of fully-random actions to seed the replay buffer / Q-table before epsilon-greedy learning begins), then evaluated over held-out evaluation seeds disjoint from training. Two baseline opponents were used throughout Levels 2-5 wherever a fixed opponent was needed: a **Random Agent (uniform random action)** and a **Tracking Agent** (a scripted paddle that follows the ball's y-position with configurable noise) — both provided by the course.

3. Level 1 — Tabular control, solo vs wall

Setting. Solo paddle practice against a wall (no opponent to train against); tabular 7-dimensional state

```
(my_paddle_y_bin, other_paddle_y_bin, ball_x_bin, ball_y_bin, ball_vx_sign, ball_vy_sign, speed_bin) .
```

Method. We implemented a single `TabularControlAgent` class supporting three TD control rules that only differ in their bootstrap target: Q-learning (`max` over next-state action values), SARSA (value of the action actually taken next), and Expected SARSA (expectation under the current epsilon-greedy policy). This isolates the effect of the TD target while holding the state representation, reward, and training budget identical across methods — a controlled comparison rather than three independent implementations.

Two experiments. We ran this comparison twice. The first pass

(`Level_1_Qlearning_SARSA_ExpectedSARSA_Comparison.ipynb`) used

`epsilon_min=0.20` and no reward shaping. Reviewing the results, we judged

the forced-exploration warm-up and the epsilon floor to be masking each

method's true converged behaviour, so we re-ran with an improved diagnostic

setup (`Level_1_Comparison_Exploration_and_Diagnostics.ipynb`): a longer

forced-exploration warm-up (400 episodes of fully random actions before

learning begins) and a lower epsilon floor (0.05), plus an optional ablation

that masks the -0.001 movement penalty out of the *training* TD target only

(official reward is still recorded separately for evaluation) to test

whether the movement penalty was discouraging useful paddle movement during

learning.

Hyperparameters. `alpha=0.1`, `gamma=0.99`, `epsilon_start=1.0`,

`epsilon_decay=0.9999`, 2000 training episodes, training seed 0, evaluated with

epsilon forced to 0 over 200 held-out episodes (seed 10,000).

Results (second, improved experiment):

Agent	mean reward	std reward	mean hits	mean steps	learned states
Q-learning	-0.0445	0.7053	2.57	437.0	33,082
SARSA	-0.0864	0.3870	1.81	336.4	32,872
Expected SARSA	-0.1202	0.1878	2.72	465.7	32,336

All three methods had a 100% miss rate against the wall (expected — Level 1

has no opponent to beat, only survival/hit count to optimize) and 0%

timeout rate. Expected SARSA had the lowest reward variance (std 0.19 vs

0.71 for Q-learning), consistent with its theoretically lower-variance

bootstrap target, but Q-learning had the highest mean reward of the three in

this run.

Failure case. Replay diagnostics (seed 42) show STAY as the single most

common action for all three agents even though it is never a majority of

steps: Q-learning 85/190 STAY (UP 50 + DOWN 55 = 105 movement actions),

SARSA 110/300 STAY (UP 98 + DOWN 92 = 190 movement actions), Expected SARSA 117/300 STAY (UP 93 + DOWN 90 = 183 movement actions). All three lost or timed out against the wall in this single-seed replay (Q-learning: right_scored after 190 steps; SARSA and Expected SARSA: timeout after the full 300-step replay budget, both credited to "right" on cumulative reward) — a reminder that "learned" here still means a policy far from solving even the no-opponent Level 1 task after 2000 training episodes.

Limitation we flag explicitly: all Level 1 numbers above come from a **single training seed**. We did not have time to repeat with 3-5 independent seeds and report mean±std, which both Level 1 notebooks themselves note as a requirement for a statistically solid conclusion. This is the single largest methodological gap in our submission — see Section 10.

4. Level 2 — Tabular control vs a scripted opponent

Setting. Left agent against a scripted opponent, tabular state (same 7-dim representation as Level 1).

4.1 Algorithm comparison

We first compared Q-learning vs SARSA with the original, unshaped reward (`Level_2_Qlearning_vs_SARSA_Normal.ipynb`), same hyperparameters as Level 1 (alpha=0.1, gamma=0.99, epsilon_decay=0.9999, 2000 episodes, 400-episode forced-exploration warm-up), evaluated over 200 held-out episodes (seed 10,000) against the same opponent used in training, plus the Random and Tracking baselines:

Agent	mean reward	win rate	mean hits	mean steps
Q-learning	2.1649	0.2150	1.11	191.5
SARSA	1.8624	0.1850	0.94	170.8
Random Agent	1.5015	0.1550	0.77	141.8
Tracking Agent	7.6266	0.7550	7.98	741.8

Both learned agents clear the Random baseline on every metric, but both are

far below the scripted Tracking Agent, which simply tracks the ball's `y` position. Q-learning outperformed SARSA on every metric in this run, so we fixed the algorithm to Q-learning for the exploration-schedule study below.

4.2 Exploration schedule comparison

With the algorithm fixed to Q-learning, `Level_2_Qlearning_Exploration_Schedules.ipynb` compares three epsilon schedules over the same 1600 post-warm-up episodes:

Constant (epsilon=0.20 throughout), **Linear decay** (1.00→0.05), and

Exponential decay (1.00→0.05):

Schedule	mean reward	win rate	mean hits	mean steps
Constant	1.6729	0.1650	1.07	191.1
Linear	2.0746	0.2050	1.245	224.6
Exponential	1.8778	0.1850	1.025	186.9

Linear decay produced the best win rate, reward, and hit count of the three schedules on this evaluation set. Our reading: constant epsilon keeps injecting 20% random actions even late in training, directly costing win rate at evaluation-adjacent behaviour; the two decay schedules let the policy exploit more as training progresses, with linear's more gradual reduction giving slightly more total learning signal than exponential's fast initial drop in this particular run.

Failure case. One exponential-decay replay (seed 42) shows the agent winning quickly (138 steps, hits=1) purely by early positioning luck rather than sustained rally play — a reminder that a single replay is not representative and reinforces the need for the multi-seed evaluation we flag as a limitation.

5. Level 3 — Robustness across an opponent zoo

Setting. Left agent evaluated against an "opponent zoo" (random, weak, tracking, strong, and a mixed "zoo" opponent that samples among them), tabular state.

Method. We trained a single Expected SARSA agent

(`Level_3_Expected_SARSA_Robustness.ipynb`) directly against the mixed

`training_opponent="zoo"` setting ($\alpha=0.1$, $\gamma=0.99$, $\epsilon_{\text{decay}}=0.9999$,

2000 episodes, 400-episode warm-up), then evaluated the **same** trained

agent separately against each individual opponent in the zoo (200 episodes

each, seed 10,000), rather than training five separate specialists — directly

testing generalization, per the level's research question.

Per-opponent results:

Opponent	mean reward	win rate	mean hits	mean steps
random	2.4932	0.2500	0.60	115.0
weak	1.5004	0.1500	0.88	163.1
tracking	0.8090	0.0800	1.13	195.7
strong	0.6597	0.0650	1.20	206.9
zoo (mixed)	1.7062	0.1700	0.87	158.1

Robustness summary: average win rate 0.143, worst-case win rate 0.065

(against `strong`), win-rate range 0.185. The **strong** opponent is the

hardest by both win rate and reward — expected, since it is presumably the

best-tuned scripted player in the zoo.

Failure case. The replay against `strong` (seed 11) is a clean, fast

loss: only 3 DOWN and 32 STAY actions over 35 steps (hits=0, reward=-0.003) —

the agent barely reacts before conceding the point, a 21px paddle-movement

range versus roughly 3x that for the wins against weaker opponents in the

same replay set. This suggests the single zoo-trained policy has not learned

a defense fast/precise enough for the strongest scripted opponent, even

though it generalizes reasonably (0.15-0.25 win rate) against the weaker

members of the zoo.

6. Level 4 — Pixel-only observation, two architectures

Setting. Pixel-only observation against scripted opponents (image

observation, no tabular state available). We built and compared two independent agents, both against the `weak` scripted opponent, both using standard DQN infrastructure (replay buffer, target network, Huber loss, Adam, gradient clipping at `max_norm=10.0`) — the only planned difference is the observation representation, isolating that variable.

6.1 Feature-DQN

Rather than a CNN, `Level_4_Feature_DQN_Baseline.ipynb` first extracts approximate (x, y) coordinates of both paddles and the ball from the raw RGB frame plus their frame-to-frame deltas (12-dim feature vector), fed to a small MLP (`Linear(12→128)→ReLU→Linear(128→128)→ReLU→Linear(128→3)`). This was deliberately built and validated *before* attempting a CNN, to confirm the two-agent DQN training loop itself was correct on a cheaper representation first.

Hyperparameters: `lr=1e-4`, `gamma=0.99`, `epsilon_decay=0.9999`, buffer capacity 50,000, `batch_size=64`, `min_buffer_size=2,000`, `target_update_steps=1,000`, `train_frequency=4`; 2,000 training episodes, 400-episode warm-up. Trained on an NVIDIA RTX 5050 laptop GPU in 2441.9s (305,564 environment steps, 75,892 optimization steps).

Evaluation (200 episodes, seed 10,000, vs weak):

Agent	mean reward	win rate	mean hits	mean steps
Feature-DQN	1.9836	0.2050	1.075	172.2
Random Agent	1.5015	0.1550	0.77	141.8
Tracking Agent	7.6266	0.7550	7.98	741.8

Feature-DQN modestly beats Random but sits well below the scripted Tracking Agent — comparable in shape to the Level 2 tabular results, which makes sense since both are learning from essentially the same underlying positional information (engineered coordinates vs binned tabular state).

Failure case. Replay (seed 22): a loss with heavy erratic movement (107

UP vs 24 DOWN vs 16 STAY, 168px paddle range) yet only 1 hit before conceding — movement without control, suggesting the learned policy reacts to noise in the extracted-coordinate features rather than a stable tracking strategy.

6.2 CNN-DQN

`Level_4_CNN_DQN_Pixel_Baseline.ipynb` keeps every DQN infrastructure choice

identical to the Feature-DQN baseline and only swaps the observation

representation: raw pixels, grayscaled, resized to 84×84, stacked 4 deep,

through a standard Atari-style conv stack

(`Conv(4→32, k8, s4)→Conv(32→64, k4, s2)→Conv(64→64, k3, s1)→Linear(3136→512)→Linear(512→3)`).

Buffer capacity was reduced to 30,000 (vs Feature-DQN's 50,000) because CNN

optimization is more memory/compute expensive per step; epsilon decay was

slowed slightly (0.99995 vs 0.9999) to compensate for the harder

representation-learning problem.

Training took 6992.4s (~1h56m) — **2.9x longer** than Feature-DQN's 2441.9s

for the same 2,000 episodes, confirming the expected extra cost of learning

perception and control jointly from pixels.

Evaluation (200 episodes, seed 10,000, vs weak):

Agent	mean reward	win rate	mean hits	mean steps
CNN-DQN	6.9040	0.7000	3.67	482.4
Feature-DQN	1.9836	0.2050	1.075	172.2
Random Agent	1.5015	0.1550	0.77	141.8
Tracking Agent	7.6266	0.7550	7.98	741.8

CNN-DQN's win rate (0.70) approaches the scripted Tracking Agent (0.755) and

clears Feature-DQN by a wide margin (0.20). The training curve shows *why*:

reward/win-rate stayed modest through episode ~800 (comparable to

Feature-DQN's range) and only broke out sharply after episode ~1300 (reward

2.93→6.17→6.99 between episodes 1300-1900) — consistent with the notebook's

own hypothesis that pixel learning needs more episodes because the network

must learn to *see* the game before it can learn to *play* it, whereas

Feature-DQN starts with the positional information already extracted.

Failure case. Where Feature-DQN's failure replay showed erratic movement, CNN-DQN's equivalent replay (seed 22) is a **win** (44 UP / 32 DOWN / 62 STAY, 1 hit, reward +9.98) — the qualitative gap between the two architectures is visible even in a single fixed-seed replay, not just aggregate statistics.

7. Level 5 — Two learned agents, self-play

Setting. Left and Right both learned agents, pixel/learned-agent interface, trained and evaluated together. This is the level Duan Xu took furthest, working with Claude Code as a coding assistant and building directly on Manan Bimal Mehta's Level 4 Feature-DQN and CNN-DQN checkpoints.

Design. Both architectures follow the same two-stage design, split to mirror Level 4's 4A/4B structure:

- **5A — fixed-opponent baseline.** A fresh learner (left) trains against the *frozen* Level 4 model of matching architecture (right). Only the left agent's transitions are optimized; this validates the two-agent `step_due1` training loop and produces a stable starting policy.
- **5B — self-play with an opponent pool, warm-started from 5A.** The 5A checkpoint continues training. The right-side opponent per episode is, with probability 0.8, a **live mirror** of the agent's own current weights (true self-play), and with probability 0.2, a **frozen snapshot** sampled from an 8-slot opponent pool (FIFO-evicted, slot 0 permanently pinned to the same frozen Level 4 model used in 5A, later slots populated by snapshotting the agent's own weights every 200 episodes). This mirrors the Level 3 opponent-zoo idea specifically to guard against the classic self-play failure mode where two live agents chase each other's latest exploit in a cycle rather than converging to a robust policy. On warm start, epsilon reopens modestly (0.3, not reset to 1.0) since the existing

5A policy is not garbage — only the opponent distribution is changing.

Both stages, both architectures: 1500 episodes, `level=5`, `render_scale=2`,

`max_seconds=15` during training (the official 60s would be far too slow for

1500-episode runs — evaluation cells use the official `max_seconds=30`).

Reproducing this at full scale takes hours (~80 min for Feature-DQN,

~160 min for CNN-DQN on an RTX 3070 Ti — see `README.md §4`), so the actual

training that produced the submitted checkpoints was run as a converted,

non-interactive Python script rather than inside the Jupyter kernel, with

console output captured to `smoketest_logs/*_FULL_TRAIN_log.txt`. The Level

5 notebooks' own cell outputs have been backfilled with this run's real

data (see the disclosure note at the top of each notebook) — no numbers in

this section or those notebooks are fabricated; where a metric was not

captured by the periodic console log (hits/steps/loss moving averages), we

say so rather than inventing a curve.

7.1 Final evaluation — 5A vs 5B, same fixed benchmarks

50-episode evaluation, `level=5`, `max_seconds=30`, against frozen Level 4,

Random, and Tracking (Tracking placed on the left side — it is side-aware

and reads unmirrored `info`, so this placement is required for a valid

comparison, see `README.md §7`):

Feature-DQN

Opponent	5A win rate	5B (self-play) win rate
Frozen L4	0.54	0.24
Random	0.48	0.28
Tracking (right)	0.02	0.02
5B vs 5A head-to-head	—	5B wins 44%
5B self-mirror	—	20% / 80% split (non-degenerate)


CNN-DQN


Opponent	5A win rate	5B (self-play) win rate
Frozen L4	0.20	0.72

Random	0.46	0.98
Tracking (right)	0.02	0.20
5B vs 5A head-to-head	—	5B wins 86%
5B self-mirror	—	42% / 58% split (non-degenerate)



7.2 Diagnosis: self-play helped CNN-DQN and hurt Feature-DQN

The two architectures show **opposite** self-play outcomes, and the training curves explain why (Figures 1-4, reconstructed from the real 16-point periodic training log — reward and win rate against whichever opponent was sampled that 100-episode window):





Feature-DQN 5B's win rate against its own training-time opponents peaks at 81% around episode 500, then declines almost monotonically to 27% by episode 1400, with a small partial recovery to 30% by episode 1500 where training stopped. Reward tracks the same arc (peak ≈ 8.0 at episode 500 $\rightarrow$ trough ≈ 2.7 at episode 1400). Training was cut off mid-decline, not at a stable endpoint, which directly explains why the final 5B checkpoint underperforms both the frozen L4 baseline (24% vs 5A's 54%) and its own 5A checkpoint in a head-to-head (44%). This is the classic self-play non-stationarity failure mode the opponent pool was designed to dampen: the pool visibly slowed full collapse (the self-mirror match stayed a non-degenerate 20/80 split rather than degenerating to 0-0 draws, so the policy did not go fully inert) but did not prevent the overall decline against opponents outside the training loop.



CNN-DQN 5B, by contrast, dips through a noisy middle third (episodes 500-800, as the opponent pool fills with progressively tougher snapshots of itself) then climbs steadily through the back half to 77% by episode 1500 — still rising when training stopped, with no sign of the Feature-DQN pattern. Our reading: the much larger CNN capacity (a full conv stack versus a 12-dimensional feature MLP) gives it more room to represent a genuinely

robust policy instead of overfitting to whichever opponent it saw most recently — the same representation-capacity story that separated the two architectures at Level 4 (Section 6) appears again here, but this time it governs *self-play stability* rather than just raw performance against a fixed scripted opponent.

Cross-architecture note: the Level 4 gap (CNN-DQN clearly beating Feature-DQN) persists and widens at the 5B (final deliverable) stage — CNN-DQN 5B beats frozen L4 72% of the time vs Feature-DQN 5B's 24%. But the relationship is not simply "CNN is always better": Feature-DQN's 5A *already* beat its frozen L4 opponent more convincingly than CNN-DQN's 5A did (54% vs 20%) before self-play was introduced. Self-play was necessary for CNN-DQN to reach a strong final policy, but actively regressed Feature-DQN's policy in this run.

****Against Tracking, all four Level 5 agents are weak — this is the one opponent none of them solved.**** Tracking is the strongest scripted baseline used anywhere in this project (see Sections 4 and 6), and both Feature-DQN variants win only 2% of the time against it (5A and 5B identically). CNN-DQN 5A also wins only 2%; self-play lifts CNN-DQN 5B to 20% — a real four-fold improvement over 5A, consistent with the mean-score gap narrowing (mean score for our agent rises from 0.09 at 5A to 1.88 at 5B, while Tracking's own mean score falls from 9.39 to 6.61) — but CNN-DQN 5B still ****loses the large majority of rallies against Tracking****. So the correct reading of the self-play result is narrower than "self-play closes the gap to strong scripted play": self-play helped substantially against the frozen L4 model and Random, and helped some against Tracking, but did not make CNN-DQN 5B competitive with Tracking specifically.

7.3 Replay evidence

A recorded rally (seed 22, 5B left vs frozen L4 right, `max_seconds=15`) for

each architecture, generated from the actual saved final checkpoints:

- **Feature-DQN 5B vs frozen L4:** left action counts `{UP: 92, DOWN: 38, STAY: 74}`, right `{STAY: 204}` — the frozen L4 opponent barely needs to move; Feature-DQN 5B loses (reward -0.07 vs +10.06), consistent with the aggregate 24% win rate above.

- **CNN-DQN 5B vs frozen L4:** left `{UP: 96, DOWN: 68, STAY: 36}`, right `{UP: 77, DOWN: 79, STAY: 44}` — an actively contested rally, and CNN-DQN 5B **wins** (reward +9.96 vs -0.10).

Full interactive replays (matplotlib animation) are embedded in both Level 5 notebooks' final cells.

8. Hyperparameter tuning summary

Level	What was compared	Winner (this run)
1	Q-learning vs SARSA vs Expected SARSA (TD target)	Q-learning (highest mean reward); Expected SARSA lowest variance
1	Movement-penalty masking in TD target (ablation)	Diagnostic tool, not a final choice — flagged for further ablation
2	Q-learning vs SARSA vs scripted opponent	Q-learning
2	Constant vs linear vs exponential epsilon decay	Linear decay
4	Feature-DQN vs CNN-DQN buffer capacity (50k vs 30k), epsilon decay (0.9999 vs 0.99995)	Architecture-specific tuning, not a head-to-head sweep
5	Fixed-opponent (5A) vs self-play + opponent pool (5B)	Architecture-dependent: CNN-DQN benefits, Feature-DQN regresses
5	Opponent pool live-mirror fraction (0.8) vs pool-only	Not swept — 0.8 was fixed a priori from the Level 3 "opponent zoo" precedent

We did not have time to run a systematic learning-rate or discount-factor sweep at any level; the comparisons above (TD target, exploration schedule, observation representation, self-play vs fixed-opponent) are the hyperparameter/design axes we actually controlled for. This is discussed further as a limitation in Section 10.

9. Which levels our final agents support

Level	Attempted	Final agent(s)	Checkpoint(s)
1	Yes	Q-learning, SARSA, Expected SARSA (tabular)	q_learning.pkl , sarsa.pkl , expected_sarsa.pkl
2	Yes	Q-learning (best schedule: linear decay)	q_learning_level2.pkl + schedule variants
3	Yes	Expected SARSA trained on opponent zoo	expected_sarsa_level3_zoo.pkl
4	Yes	Feature-DQN and CNN-DQN, both vs scripted weak opponent	feature_dqn_level4.pt , cnn_dqn_level4.pt
5	Yes	Feature-DQN and CNN-DQN, both 5A (fixed-opponent) and 5B (self-play) variants	feature_dqn_level5_5a.pt / _selfplay.pt , cnn_dqn_level5_5a.pt / _selfplay.pt

Performance trend across levels: win rate against the best available

fixed benchmark (Tracking Agent where used, frozen Level 4 opponent at Level 5) improves substantially from tabular (Level 1-3, generally <25% win rate against non-trivial opponents) to CNN-DQN pixel learning (Level 4: 70% vs weak; Level 5 CNN-DQN 5B: 72% vs frozen L4). Feature-DQN is the exception to a clean monotonic story: it performs respectably as a fixed-opponent learner (Level 4: 20.5% vs weak; Level 5A: 54% vs frozen L4) but does not benefit from — and in our Level 5 self-play run, was actively hurt by — the harder self-play setting, unlike CNN-DQN.

Recommended final submission for the competition interface:

cnn_dqn_level5_selfplay.pt — the strongest, most self-play-stable agent we trained against the frozen Level 4 model and Random (72% and 98% win rate respectively), valid for either the Left or Right slot since the environment already mirrors observations for the right side. Caveat: it still loses the majority of rallies against the Tracking scripted baseline (20% win rate, Section 7.2) — "strongest agent we produced" should not be read as "solves the game," and a competition matchup against a Tracking-quality opponent (or another team's strong scripted agent) is the scenario where this checkpoint is most likely to underperform its Level-4/Random numbers.

10. Limitations and future work

- **Single training/evaluation seed at Levels 1-3.** Every tabular result in this report comes from one training seed and one evaluation seed set. Both Level 1 notebooks explicitly flag this as insufficient for a strong scientific conclusion and recommend 3-5 independent seeds with $\text{mean} \pm \text{std}$ reporting. Given more time, this is the single highest-value addition to the tabular sections.
- **No systematic hyperparameter sweep.** We compared *architectural* and *algorithmic* choices (TD target, exploration schedule, observation representation, self-play vs fixed opponent) but did not sweep learning rate, discount factor, or network width/depth independently at any level.
- **Feature-DQN self-play collapse was not re-run to convergence.** Section 7.2's diagnosis is strong (a clear peak-then-decline pattern from real logged data) but we stopped at the pre-planned 1500 episodes rather than extending training to see whether the oscillation is a longer cycle that would eventually recover, or a genuine instability specific to the smaller Feature-DQN's capacity. This is a concrete, well-motivated next experiment.
- **hits/steps/loss training curves for Level 5 are unavailable.** The full 1500-episode-resolution moving averages for these three metrics were tracked in memory during training but the periodic console log (the only artifact that survived the multi-hour background run) only captured reward/win_rate/epsilon every 100 episodes. Re-running with the history dict serialized to disk periodically would fix this for future work.
- **Competition-readiness untested end-to-end.** We have not run our final agents through the teaching team's actual competition notebook interface, only through our own `evaluate_duel_agents / play_rally` calls.

11. Conclusion

Across five levels of increasing difficulty, the clearest pattern in our results is that **representation capacity governs not just raw performance but also training-dynamics stability**. At Level 4, CNN-DQN needed roughly 3x the training time of Feature-DQN but ultimately learned a much stronger policy from the same reward signal (70% vs 20.5% win rate against a fixed scripted opponent). At Level 5, that same capacity gap resurfaces as a *stability* story: CNN-DQN's self-play training climbed steadily to a strong final policy, while Feature-DQN's self-play training peaked early and then regressed for the remainder of training, ending up worse than its own fixed-opponent baseline. The opponent pool we built specifically to guard against self-play instability partially worked (no agent ever collapsed to degenerate 0-0 draws) but was not sufficient to prevent Feature-DQN's decline — a concrete, evidenced limitation rather than a hand-wavy one, thanks to the real per-100-episode training curves recovered from the actual training run's logs.

Figures in Section 7 and the checkpoint/reproducibility details referenced throughout this report are available in `report_assets/` , `saved_agents/` , and `README.md` at the project root.